

A person is playing a guitar in a workshop. On a wooden desk, there is a laptop displaying a video call with four participants. Next to the laptop is a smartphone and a small green pedal. In the background, there is a large grey electronic device and a desk lamp. In the foreground, a yellow and black power supply unit with the 'EIK' logo is connected to cables.

# EIK

## Eik Workshop ADC '22

Hassle-Free Embedded Development on Your Computer

# Elk Workshop ADC '22

Workshop repository: [github.com/elk-audio/adc-22.git](https://github.com/elk-audio/adc-22.git)

Docs	<a href="https://elk-audio.github.io/elk-docs"><u>elk-audio.github.io/elk-docs</u></a>
Code	<a href="https://github.com/elk-audio"><u>github.com/elk-audio</u></a>
Community Plugins	<a href="https://elk-audio.github.io/elk-community"><u>elk-audio.github.io/elk-community</u></a>
Website	<a href="https://elk.audio"><u>elk.audio</u></a>
Forum	<a href="https://forum.elk.audio"><u>forum.elk.audio</u></a>

# Outline - First part

Elk Audio Introduction (talk)

Demonstrations:

- 1 Set up a chain of plugins
- 2 Write a control application that uses SUSHI's gRPC API
- 3 Implement the control of your embedded device using
  - a Physical controls
  - b Remote GUIs
  - c End-user development tools
- 4 Use additional tools to monitor performance and problems
- 5 Q&A

# Outline - Second part

Work independently, with our assistance

Using the same tools we demonstrated

Create a prototype of an embedded audio device:

- A simple synthesiser

- Or stompbox pedal

A few Elk hardware units will be available if you want to run your experiments on real hardware

# Requirements

## Laptop

- macOS (10.15 or later)
- Linux with a recent distribution and JACK audio server

## Basic knowledge of one of these two languages:

- Python (recommended)
- C++

# Requirements

## Optional

- Small MIDI controller
- JUCE dev environment
- Elk Audio OS SDK - if you want to cross-compile



# Technology

Custom Linux Distribution  
Xenomai realtime Kernel  
ARM & x86

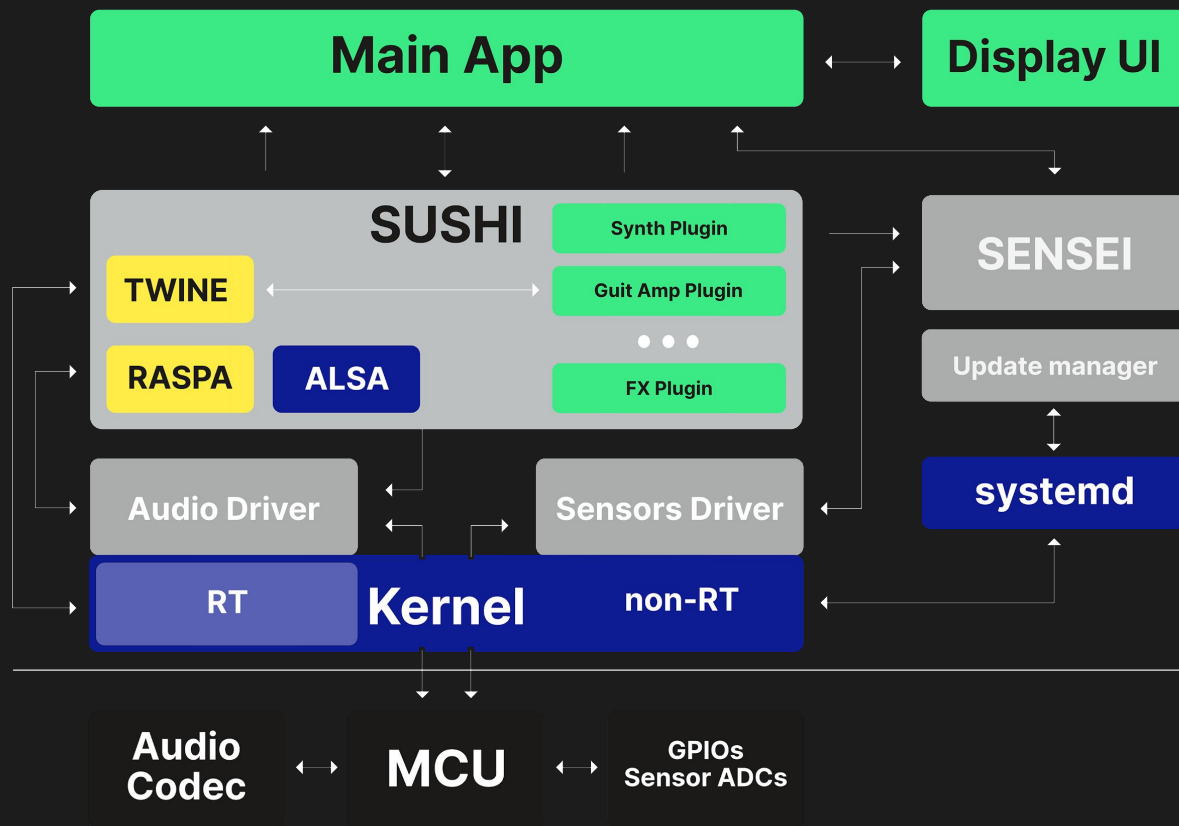
Connectivity: WiFi, BLE  
Plugin support: VST2, VST3, LV2  
CV & Gate I/O  
Ableton Link



# Architecture

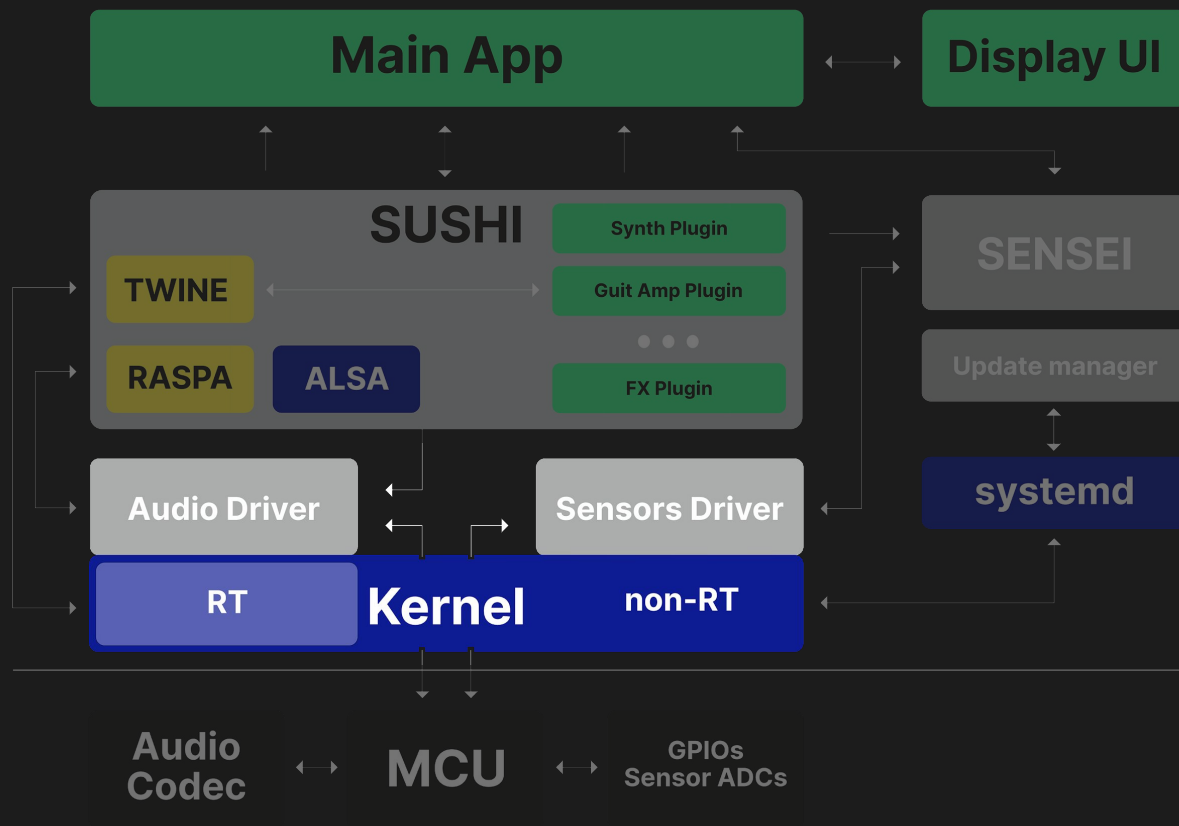


# Architecture



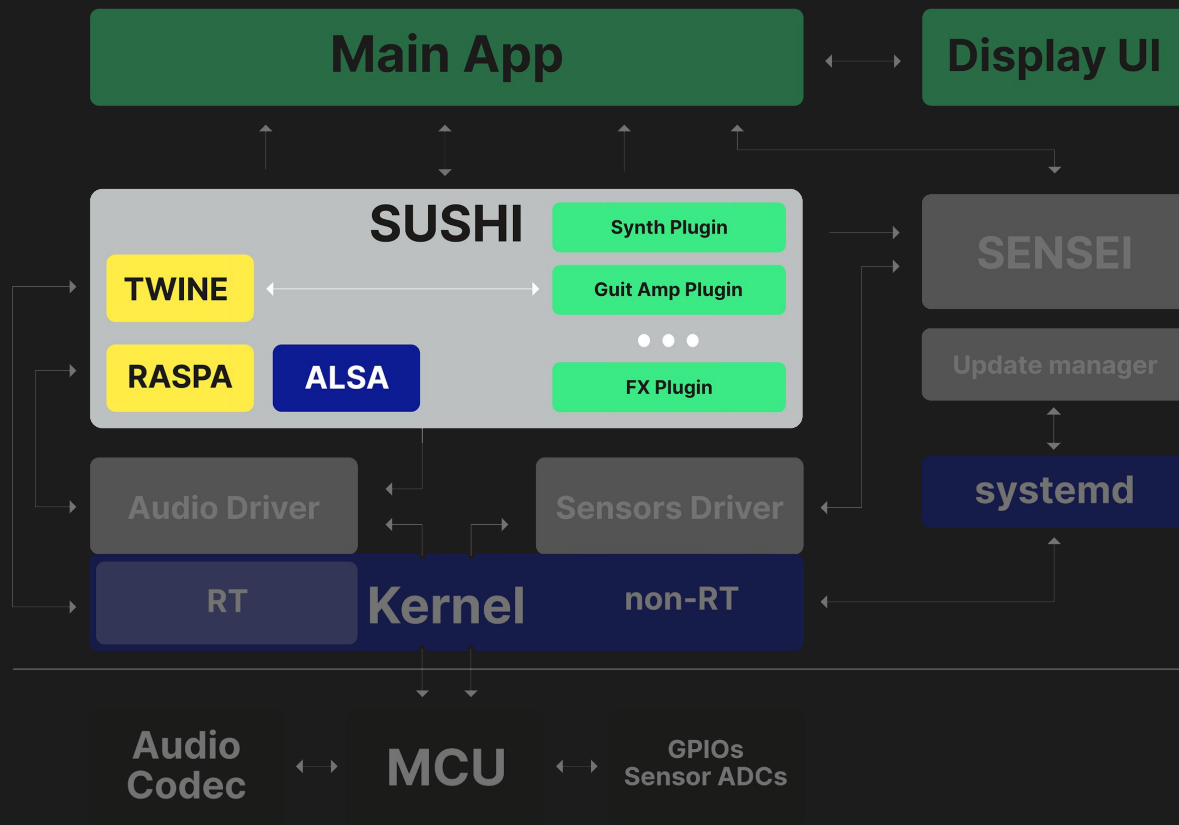
# Architecture

Dual Kernel



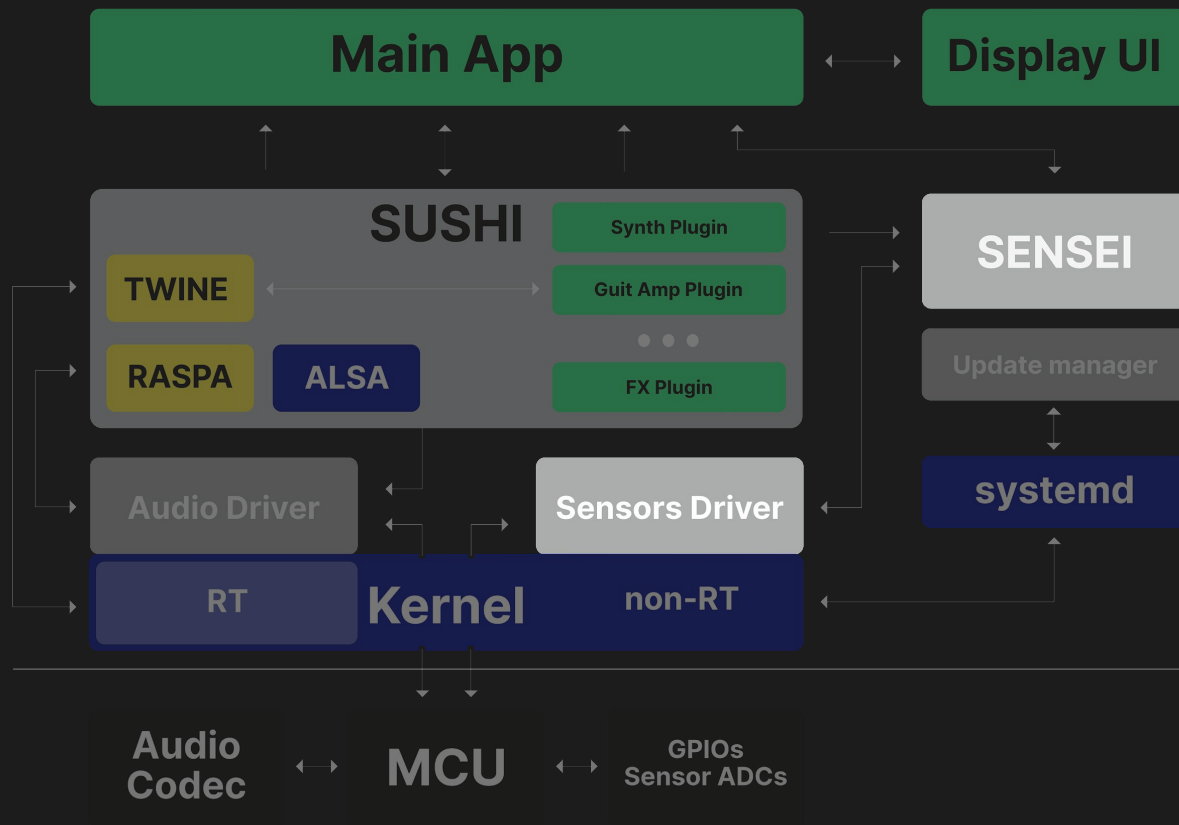
# Architecture

Dual Kernel  
Plugin Host



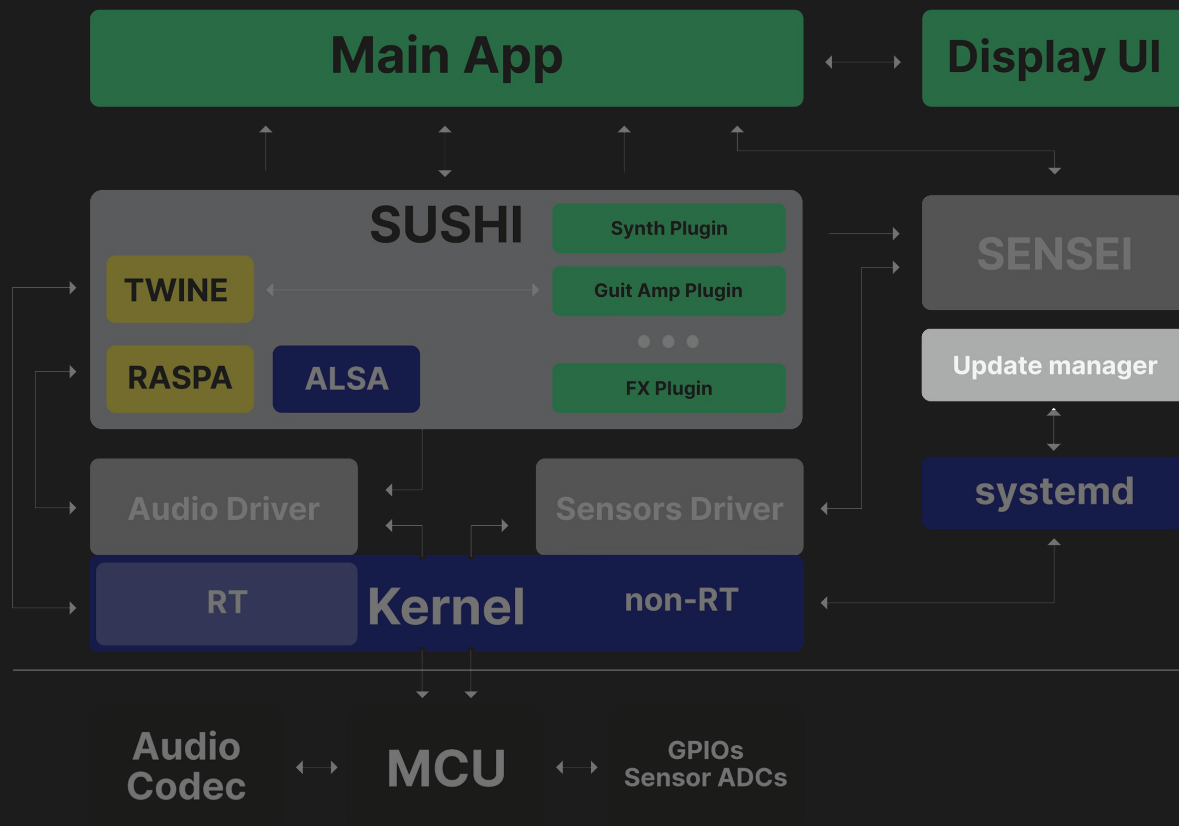
# Architecture

Dual Kernel  
Plugin Host  
Sensor Daemon



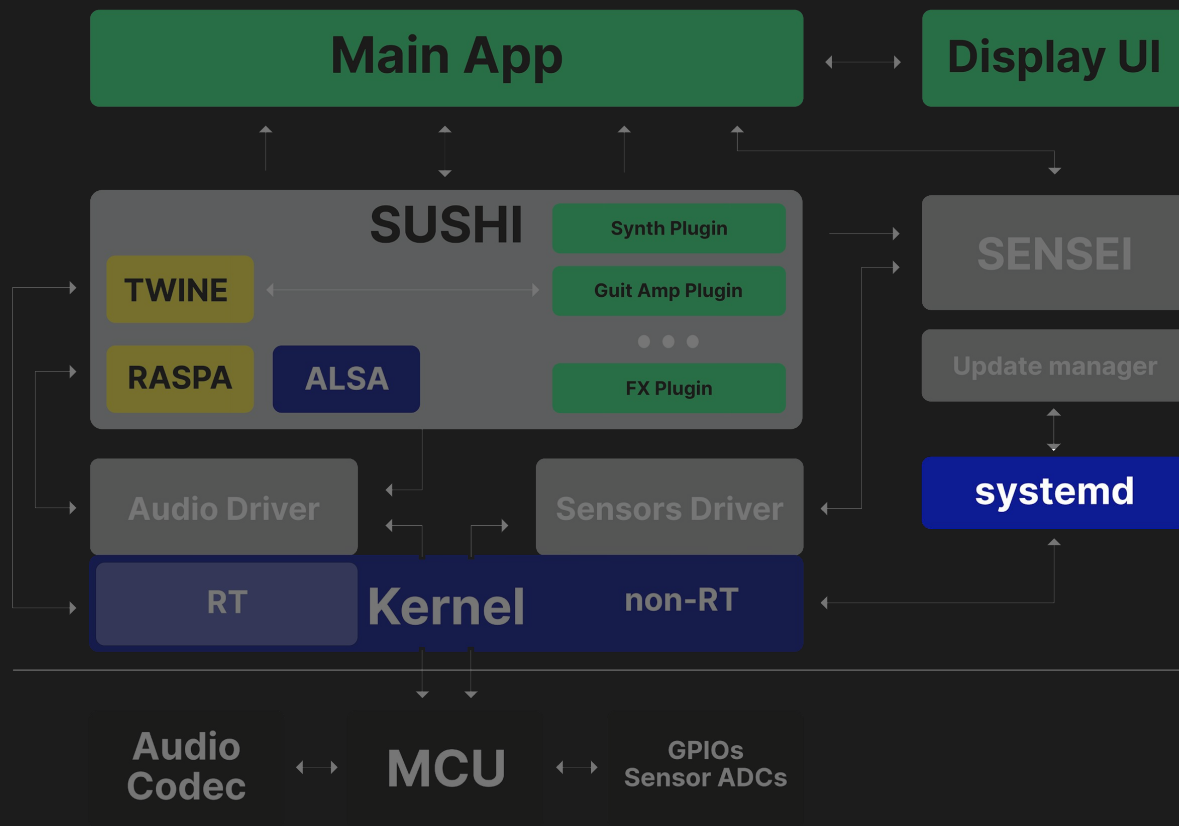
# Architecture

Dual Kernel  
Plugin Host  
Sensor Daemon  
Software Update



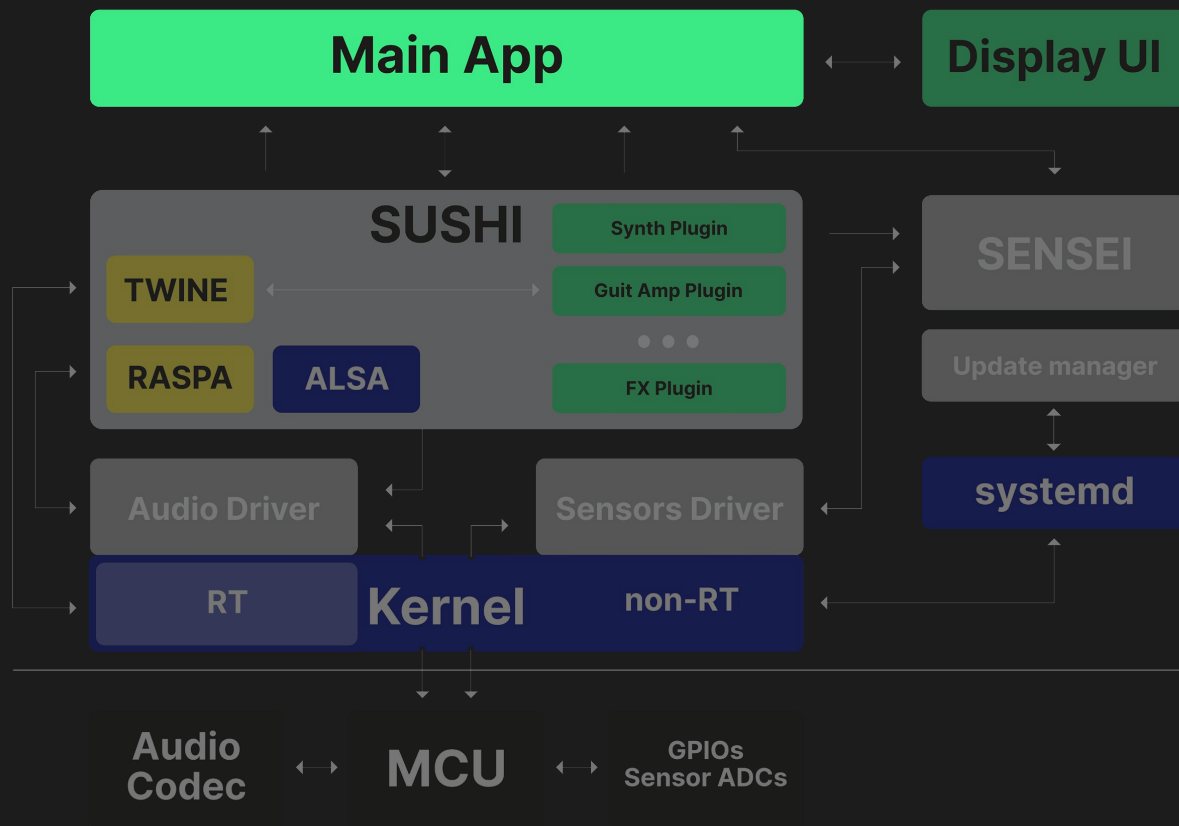
# Architecture

Dual Kernel  
Plugin Host  
Sensor Daemon  
Software Update  
Systemd



# Architecture

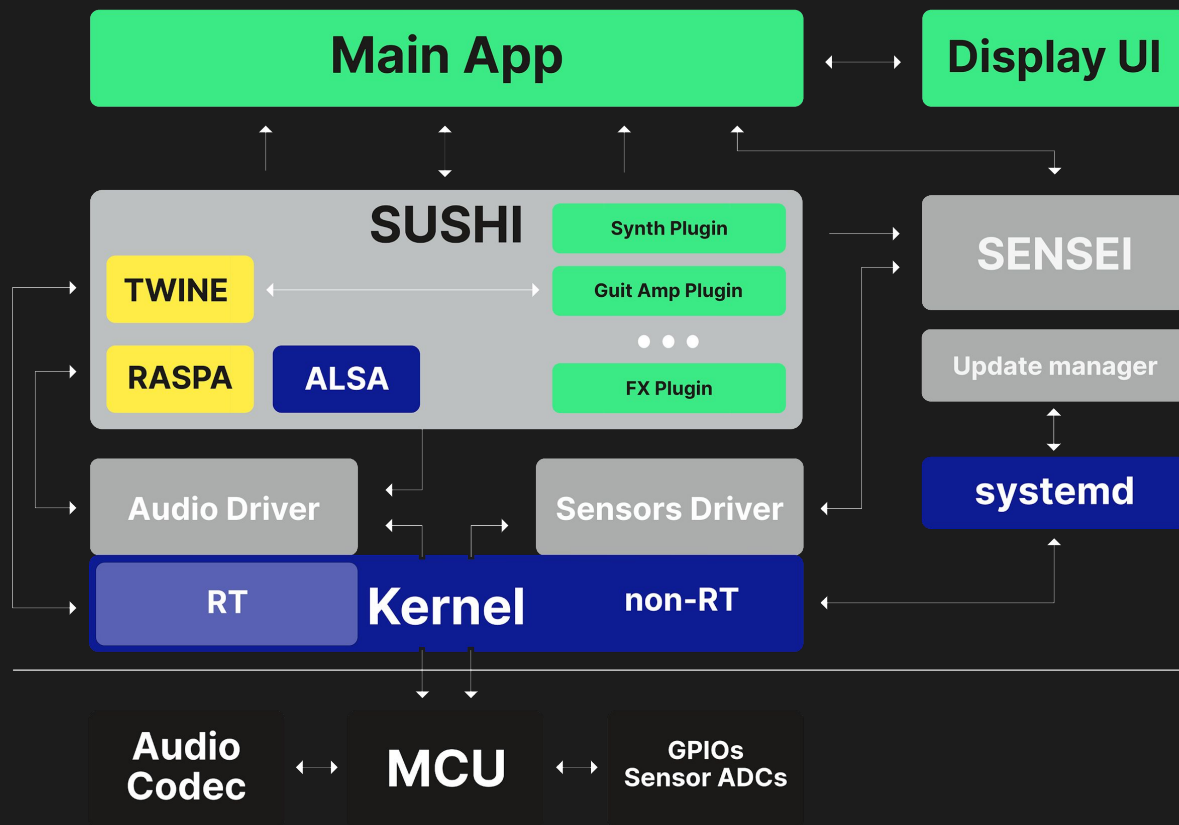
Dual Kernel  
Plugin Host  
Sensor Daemon  
Software Update  
Systemd  
Main App





# Architecture

Dual Kernel  
Plugin Host  
Sensor Daemon  
Software Update  
Systemd  
Main App



# Sushi DAW overview

Headless host with full run-time control  
MIDI, gRPC, OSC

Hosts VST 2.4, VST 3.7 and LV2 plugins


Low latency performance

Multithreaded audio processing support

Ableton Link support

Audio connections through  
Raspa, Jack, Portaudio, and file I/O

Simple scripting configuration



# Sushi DAW configuration


JSON configuration file

example coming in the next slides

gRPC

as the final exercise

<https://elk-audio.github.io/elk-docs>



**Distributed  
Systems**

**Messaging  
Patterns**



# Transport Layer

Messaging inherently depends on transport layer

UDP - TCP

Each has advantages and disadvantages

# Transport Layer

UDP - Connectionless:

- Messages are broadcasted

- Receipt not guaranteed

- No confirmation expected

- No reply

- Nor is order guaranteed

# Transport Layer

TCP – Connection Oriented:

- No Broadcasting

- Potentially high latency and jitter

- But reliable messaging

Each has its advantages and disadvantages

A conscious choice is needed!

# Message Passing

The choice is made to:

Use the advantages of connectionless messaging

...With the disadvantages this entails

It's for direct messaging, with no reply or confirmation expected - and the lowest possible latency.

Think of MIDI as an example!



# Request / Response

Client requests → server replies

Like a function invocation

With receipt confirmation

With return value

Synchronous or asynchronous

Think of a web 1.0 server

# **Publish / Subscribe**

Clients subscribe to server notifications

Servers notify on each change

Notifications stop when:

- A server is offline

- Clients unsubscribe

Think of web 2.0: push notifications

# Protocols and Libraries



# Messaging - Open Sound Control

25 years old now!

Originally for music performance data

- Shared between instruments & computers

- In widespread use today

Official spec defines socket messaging

Common for request / response, and publish / subscribe:

- Building on official spec

- Add-ons also exist, e.g. OSC Query

# OSC is widely supported

## **Control applications:**

TouchOSC - TWO - OpenStageControl - OSC/Pilot - Lemur

## **DAW's:**

Ableton Live - BitWig - Reaper - MOTU Digital Performer - LV2 Ingen

## **Plugins:**

Native Instruments' Reaktor - FAW Circle^2 ...

## **Visuals, Stage Lighting & Projection Mapping:**

Resolume - TouchDesigner - Unreal Engine - Notch - VDMX - MadMapper ...

# OSC - Messages for a synth instance

/synth/oscillator\_1/frequency, f, 440.0

/synth/filter/cutoff, f, 65.0

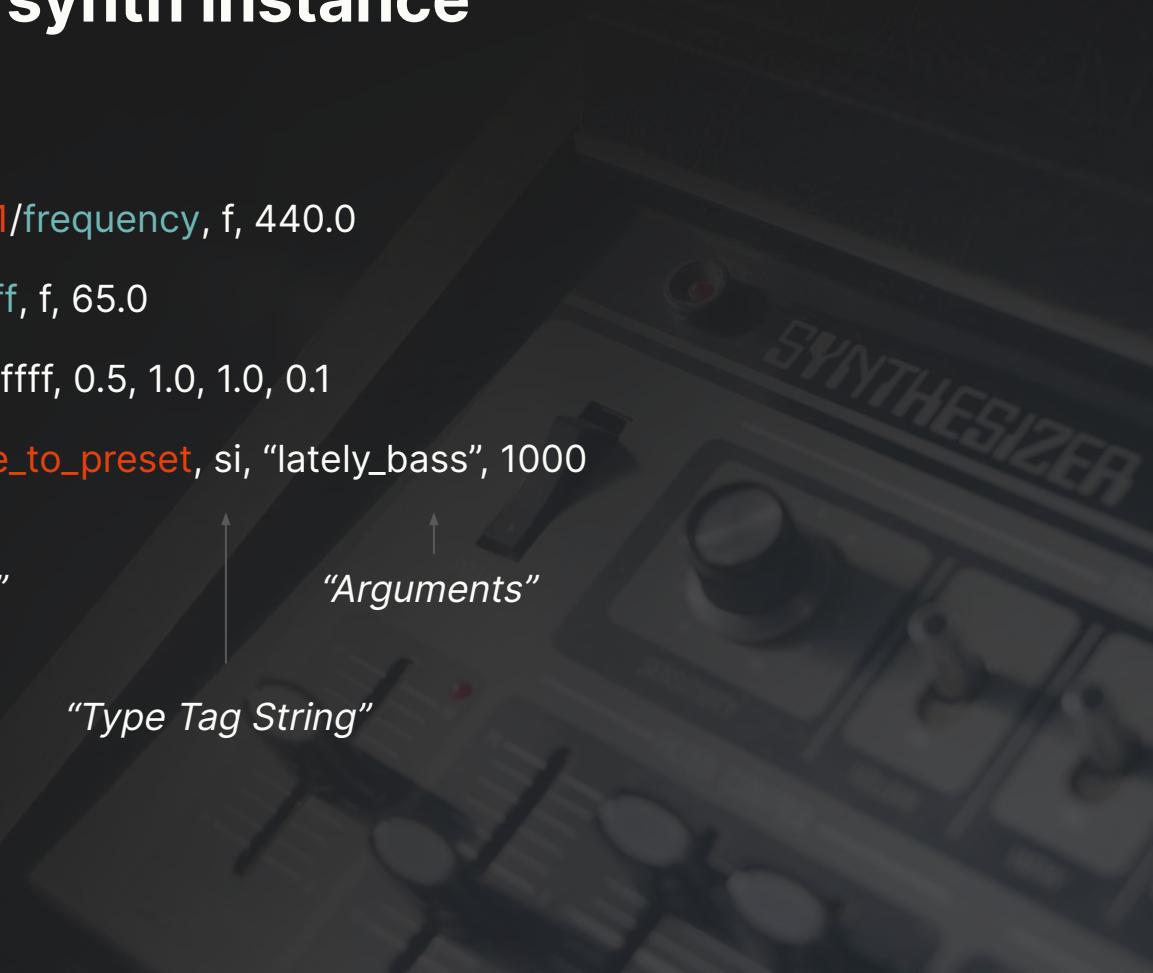
/synth/envelope, ffff, 0.5, 1.0, 1.0, 0.1

/synth/interpolate\_to\_preset, si, "lately\_bass", 1000


↑  
"Address Pattern"

↑  
"Type Tag String"

↑  
"Arguments"



# OSC - Simple synthesizer namespace



# OSC Query

Adds rich discovery features for OSC

In its core is a simple HTTP server providing:

- Existence discovery

- Namespace & State

Much optional extended functionality

LGPL C++ library available: **libossia**



# RPC Protocols

A Remote Procedure Call protocol is needed

For native request / response & publish / subscribe

Think of invoking remote methods

As if they are local

# Google Remote Procedure Call


Requires a statically defined API using:

*“Protocol Buffer Interface Description Language”*

.protobuf

Describes methods and data structures

Transcompiled for each language



# gRPC - protocol buffer file

```
service Plugin {  
    rpc SetParameter (SetRequest) returns (SetReply) {}  
}  
  
// Request message with float value:  
message SetRequest {  
    float value = 1;  
}  
  
// Response message with value as string:  
message SetReply {  
    string value = 1;  
}
```

# gRPC - In Use

Server

```
#include "api.grpc.pb.h"
```

```
Status SetParameter(ServerContext* context,  
                    const SetRequest* request,  
                    SetReply* reply) override {  
    reply->set_parameter(_to_label(request->value()));  
  
    return Status::OK;  
}
```

# gRPC - In Use

## Client

```
#include "api.grpc.pb.h"
```

```
std::string SetParameter(float value) {  
    (...)  
    request.set_parameter(value);  
    Status status = stub_->SetParameter(&context, request, &reply);  
    if (status.ok()) return reply.value();  
    else             return "RPC failed";  
}
```

# Implementation Challenges



# Concurrency



Control handling implementation in separate thread:

It cannot be in the UI or audio processes

The usual concurrency issues arise

Thankfully solutions are extensively covered!

Fabian Renn-Giles' & Dave Rowland's excellent ADC 19 talks:

"Real-Time 101" parts 1 & 2.

# Reliability



For messaging on ethernet LAN, UDP is fine.

Problems arise for critical messages

For example Note On / Off

- If message is lost or out of order

- Logic needs to be able to recover



# Feedback loops

Can arise between nodes - when both update each other of same change

Solution:

- UIs don't echo on changes from remote messages

- Only when set from UIs' controls

Complementary solution:

- For *continuous* properties, like e.g. volume

- Only send a notification if value changes

# Where to implement remote control

In the hosting DAW - or the plugin itself?

Bypassing host plugin API allows complex control

But challenges arise

Reflected parameters need to be synced between:

External control API

And host plugin API

Collisions are possible and need resolving

# Bandwidth / Rate-limiting



Despite small packets

Unsurmountable message counts are possible!

Best to rate-limit messages to 25-30Hz (UI refresh rate)

Rely on smoothing to interpolate over received values

Common also for plugin API parameters

**OSC & gRPC  
Together address  
all needs**



# Latency: OSC vs gRPC (simple messaging vs RPC)

OSC: UDP is standard

gRPC uses HTTP/2 by default

OSC is low-latency

(but stateless and connectionless)

Ideal for setting property state, and reflecting state changes

# gRPC does complex architectures well

- Request / Response model
- Connection based
- Publish / Subscribe (gRPC's streaming service)
- Synchronous and Asynchronous messaging

Meaning:

- Control a full DAW without compromise
- Discover namespaces dynamically
- Complex data structures

# End-User Development

Depends on Run-Time editability:

- Human-Readable messages
- Modifying messages, and connections

With OSC you get both

Also a big ecosystem of end-user software supporting OSC

**Messaging or RPC?**

**OSC or gRPC?**

Yes!

Combine their strengths



# Elk's Tools



# Sushi - gRPC excerpt

```
service AudioGraphController {  
    ...  
    rpc CreateTrack (CreateTrackRequest)           returns (GenericVoidValue) {}  
    rpc GetTrackProcessors (TrackIdentifier)        returns (ProcessorInfoList) {}  
    ...  
}  
  
message CreateTrackRequest {  
    string name = 1;  
    int32 channels = 2;  
}  
  
message TrackIdentifier {  
    int32 id = 1;  
}  
  
message ProcessorInfoList {  
    repeated ProcessorInfo processors = 1;  
}  
  
message ProcessorInfo {  
    int32 id = 1;  
    string label = 2;  
    string name = 3;  
    ...  
}
```

# Sushi Example GUI & elkpy

## Sushi example GUI

GUI for controlling Sushi and hosted plugins

Written in Python with QT (PySide6)

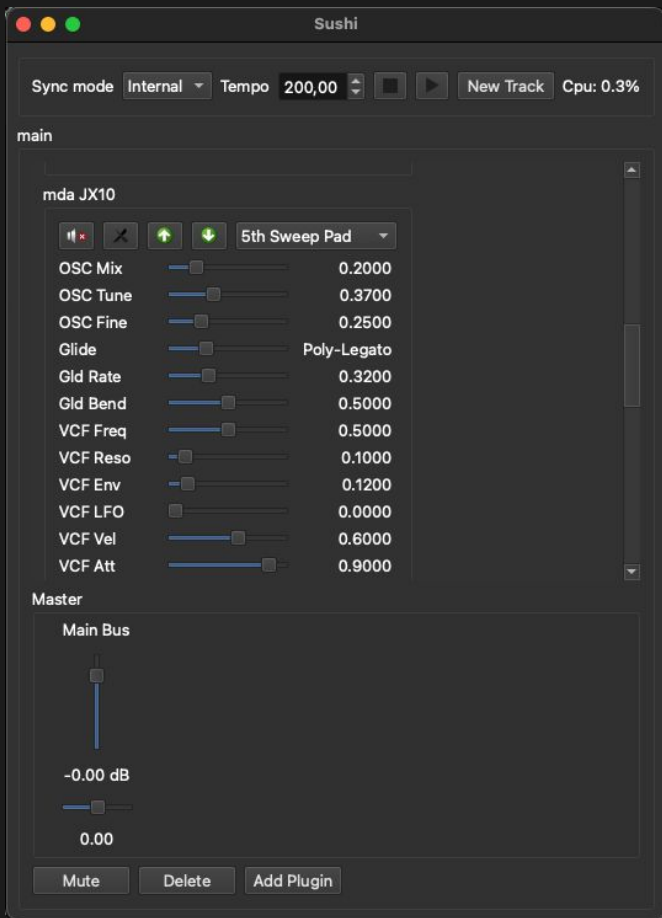
It uses the **elkpy** library:

Wrapper around Sushi's gRPC interface

Written in Python

Simplifies gRPC use

Corresponding **elkcpp** library for C++



# Sushi - elkpy example code

```
def load_plugin_on_track(self, track_id: int, plugin_spec: dict):
    path = plugin_spec['path']
    name = plugin_spec['name']
    p_type = plugin_spec['type']
    uid = plugin_spec['uid']

    try:
        self.audio_graph.create_processor_on_track(name, uid, path, p_type, track_id, 0, True)
    except Exception as e:
        print('Error loading plugin: {}'.format(e))

self.set_pan_label(txt_value)

value = self._controller.parameters.get_parameter_value(self._processor_id, gain_id)
self.set_gain_slider(value)
txt_value = self._controller.parameters.get_parameter_value_as_string(self._processor_id, self.gain_id)
self.set_gain_label(txt_value)

self._connect_signals()
```

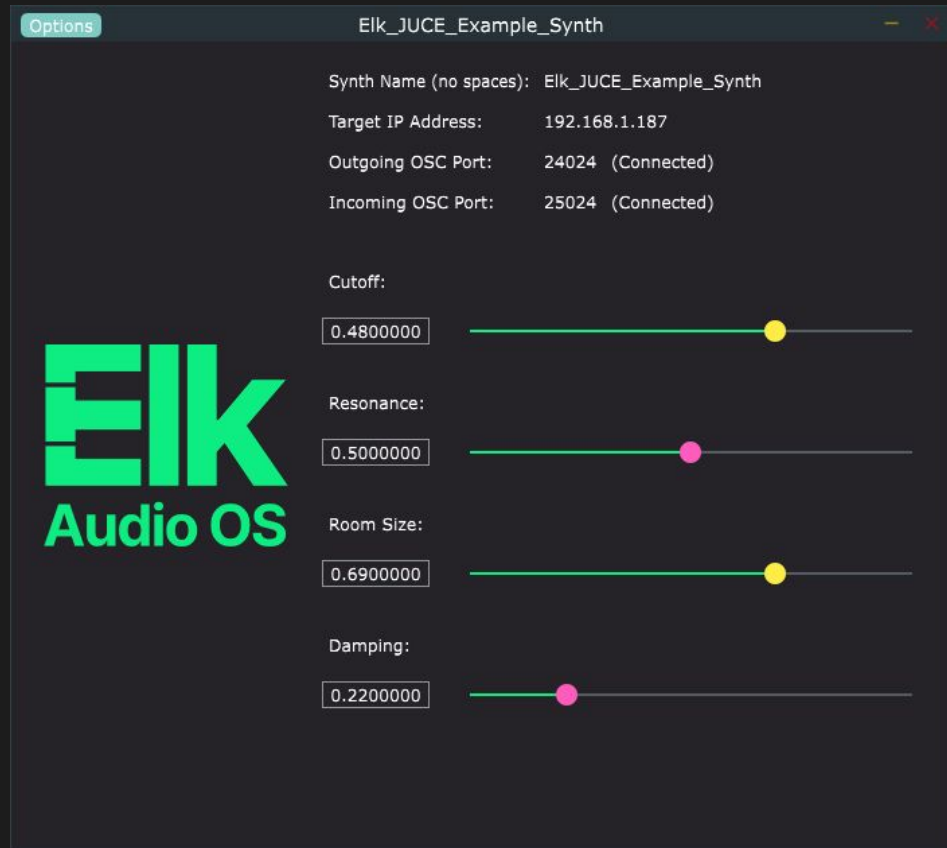
# Elk JUCE Example

Our example plugin, in JUCE

Implementing OSC control

In parallel to the plugin API

It demonstrates bridging between these



## What we've covered:

1. Contexts for headless audio
2. Distributed Systems patterns
3. Challenges and solutions to those
4. Available tools
5. Our concrete choices: OSC & gRPC

# Demonstration



# Demonstration

gRPC messaging Python example demonstrates:

- request / response
- publish / subscribe
- asynchronous communication
- non-float-parameters

Elk example plugin using OSC demonstrates:

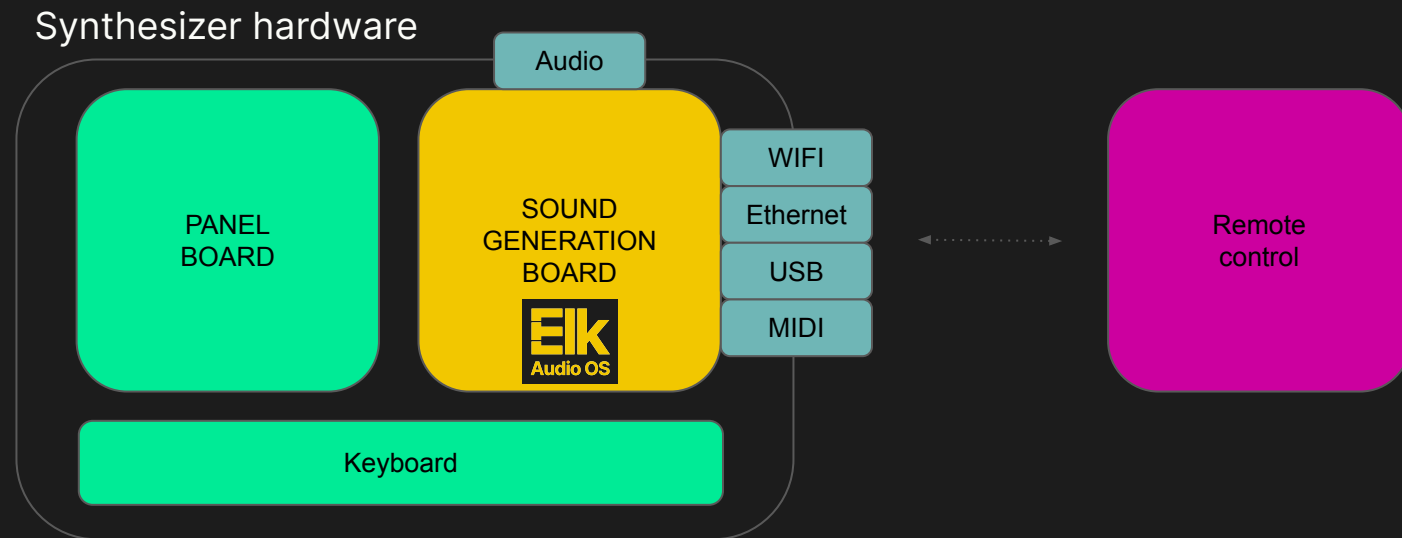
- Bypassing host plugin API
- OSC is human-readable
- Tools not made for each other can be used together, with run-time alteration
- Bridging plugin and remote API's



**Build your  
Synthesizer**



# Synthesizer platform



# Synthesizer prototyping

Developing algorithms on the target hardware

- Difficult

- Error prone


- Lack of debugging tools and connectivity

The solution

- Prototype the algorithms on your computer using sushi

- Run the same configuration on the target hardware

# Synthesizer audio processing



Example: prototype a digital synthesizer before running it on embedded hardware

# Run SUSHI

```
$ cd ~/elk-adc22/binaries
```

```
(extract binary for your platform)
```


```
$ ./sushi --portaudio -c ../config/01_synth.json
```

CTRL+C to stop

# Sushi JSON configuration example

Step 1 - the synth engine

Tracks

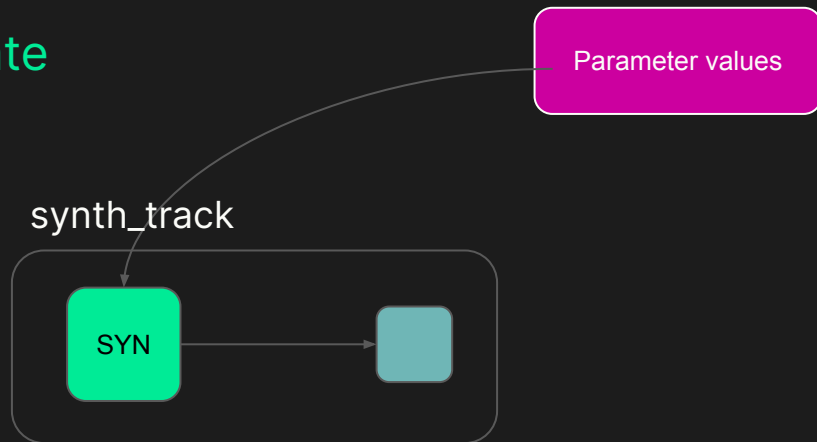


config/01\_synth.json

# Sushi JSON configuration example

Step 2 - initialize the parameters

Initial state




config/02\_initial\_state.json

# Sushi JSON configuration example

Step 3 - use a sequencer for testing...

...and use `initial_state` to set the pitch of the steps




`config/03_sequencer.json`



# Sushi JSON configuration example

Step 4 - add the insert effect

synth\_track

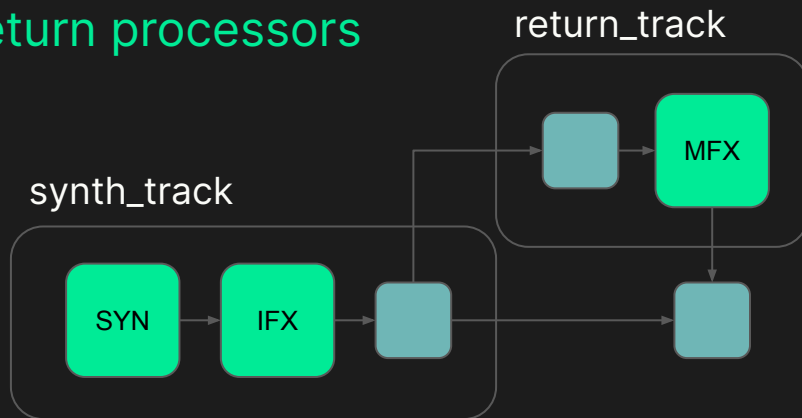


config/04\_insert\_effect.json

# Sushi JSON configuration example

Step 5 - add the master effects

Send / return processors

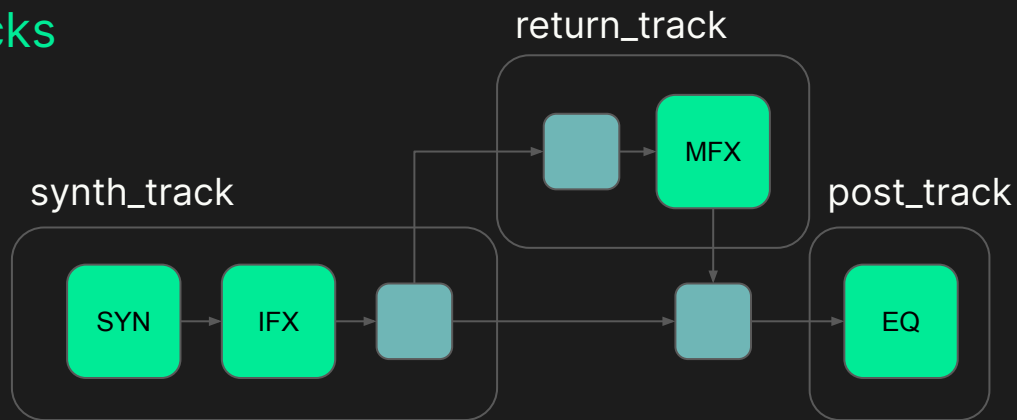


config/05\_master\_effect.json

# Sushi JSON configuration example

Step 6 - add the post EQ

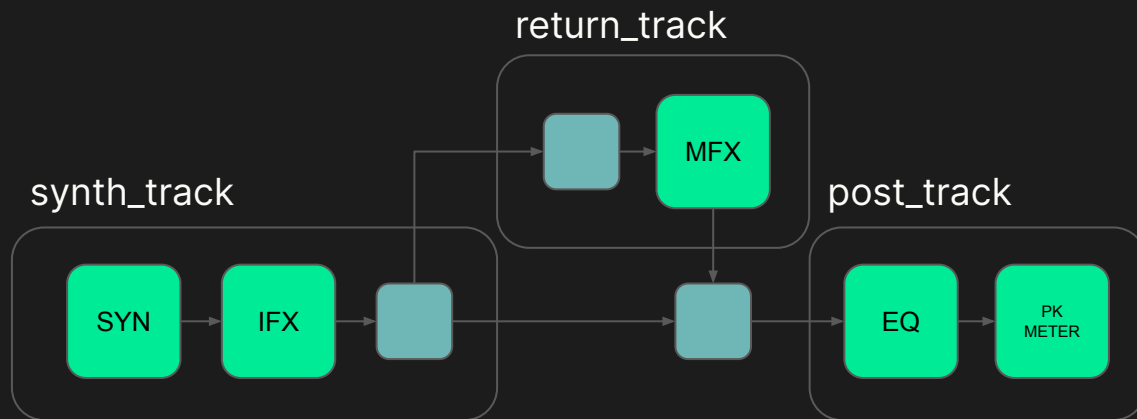
Post tracks



config/06\_post\_track.json

# Sushi JSON configuration example

Step 7 - add OSC notifications



config/07\_osc\_notifications.json

# Free exercise

Configure sushi using the gRPC API

- Create a synth chain similar to the example, but using the gRPC API
- Create some “magic buttons” to switch between different FXs / Generators

config/exercise.json



# EIK

Thank you!